

Step 1

```
import pandas as pd
```

```
import seaborn as sns
```

Step 2

```
df=pd.read_csv("/content/Churn_Modelling.csv")
```

df

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure			
	Balance	NumOfProducts	HasCrCard	IsActiveMember			EstimatedSalary	Exited			
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
	1	1	101348.88	1							
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
	0	1	112542.58	0							
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
	1	0	113931.57	1							
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0
	0		93826.63	0							
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82		
	1	1	1	79084.10	0						
...	
									
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	
	1	0	96270.64	0							
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61		
	1	1	1	101699.77	0						
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0
	1		42085.58	1							

10000 rows × 14 columns

```
df.describe()
```

RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard
	IsActiveMember	EstimatedSalary	Exited				
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
	10000.000000	10000.00000	10000.000000	10000.000000	10000.000000		

mean	5000.50000 1.530200	1.569094e+07 0.70550	650.528800 0.515100	38.921800 100090.239881	5.012800 0.203700	76485.889288
std	2886.89568 0.581654	7.193619e+04 0.45584	96.653299 0.499797	10.487806 57510.492818	2.892174 0.402769	62397.405202
min	1.00000 1.000000	1.556570e+07 0.00000	350.000000 0.000000	18.000000 11.580000	0.000000 0.000000	0.000000
25%	2500.75000 1.000000	1.562853e+07 0.00000	584.000000 0.000000	32.000000 51002.110000	3.000000 0.000000	0.000000
50%	5000.50000 1.000000	1.569074e+07 1.00000	652.000000 1.000000	37.000000 100193.915000	5.000000 0.000000	97198.540000
75%	7500.25000 2.000000	1.575323e+07 1.00000	718.000000 1.000000	44.000000 149388.247500	7.000000 0.000000	127644.240000
max	10000.00000 4.000000	1.581569e+07 1.00000	850.000000 1.000000	92.000000 199992.480000	10.000000 1.000000	250898.090000

df.isnull().any()

RowNumber False

CustomerId False

Surname False

CreditScore False

Geography False

Gender False

Age False

Tenure False

Balance False

NumOfProducts False

HasCrCard False

IsActiveMember False

EstimatedSalary False

Exited False

dtype: bool

df.isna().sum()

RowNumber 0

CustomerId 0

Surname 0

CreditScore 0

Geography 0

Gender 0

Age 0

Tenure 0

Balance 0

NumOfProducts 0

HasCrCard 0

IsActiveMember 0

EstimatedSalary 0

Exited 0

dtype: int64

df.skew()

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

"""Entry point for launching an IPython kernel.

RowNumber 0.000000

CustomerId 0.001149

CreditScore -0.071607

RowIndex	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure			
		Balance	NumOfProducts	HasCrCard	IsActiveMember		EstimatedSalary	Exited			
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
	1	1	101348.88	1							
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
	0	1	112542.58	0							
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
	1	0	113931.57	1							
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0
	0	93826.63	0								
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82		
	1	1	1	79084.10	0						
...	
									

9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2
	1	0	96270.64	0						
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	
	1	1	101699.77	0						
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1
	1	42085.58	1							0
9998	9999	15682355	Sabbatini	772	Germany		Male	42	3	
		75075.31	2	1	0	92888.52	1			
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1
	1	0	38190.78	0						

7523 rows × 14 columns

```
df.NumOfProducts.unique()
```

```
array([1, 3, 2, 4])
```

```
df.boxplot(column="Balance")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fd191cc9b50>
```

```
df=df[df.Balance < 150000]
```

```
df
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure				
	Balance	NumOfProducts	HasCrCard	IsActiveMember		EstimatedSalary	Exited				
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
	1	1	101348.88	1							
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
	0	1	112542.58	0							
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0
	0	93826.63	0								
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82		
	1	1	1	79084.10	0						

5	6	15574012	Chu	645	Spain	Male	44	8	113755.78	2
	1	0	149756.71	1						
...
								
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2
	1	0	96270.64	0						
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	
	1	1	101699.77	0						
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1
	1	42085.58	1							0
9998	9999	15682355	Sabbatini	772	Germany		Male	42	3	
		75075.31	2	1	0	92888.52	1			
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1
	1	0	38190.78	0						

6782 rows × 14 columns

```
df.boxplot(column="Age")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fd191bab050>
```

```
df=df[(df.Age <50) & (df.Age >20)]
```

```
df
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure				
	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited					
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
	1	1	101348.88	1							
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
	0	1	112542.58	0							
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0
	0	93826.63	0								

4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82		
	1	1	1	79084.10	0						
5	6	15574012	Chu	645	Spain	Male	44	8	113755.78	2	
	1	0	149756.71	1							
...
									
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	
	1	0	96270.64	0							
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61		
	1	1	1	101699.77	0						
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0
	1	42085.58	1								
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3			
		75075.31	2	1	0	92888.52	1				
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	
	1	0	38190.78	0							

5783 rows × 14 columns

```
df.boxplot(column="CreditScore")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fd191b53110>
```

```
df=df[(df.CreditScore>500) & (df.CreditScore<790)]
```

```
df
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure			
	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited				
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1
	1	1	101348.88	1						
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1
	0	1	112542.58	0						

3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0
	0	93826.63	0								
5	6	15574012	Chu	645	Spain	Male	44	8	113755.78	2	
	1	0	149756.71	1							
8	9	15792365	He	501	France	Male	44	4	142051.07	2	
	0	1	74940.50	0							
...
									
9990	9991	15798964	Nkemakonam	714	Germany			Male	33	3	
		35016.60	1	1	0	53667.08	0				
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	
	1	0	96270.64	0							
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61		
	1	1	1	101699.77	0						
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0
	1	42085.58	1								
9998	9999	15682355	Sabbatini	772	Germany			Male	42	3	
		75075.31	2	1	0	92888.52	1				

4993 rows × 14 columns

```
churn_yes=df['Tenure'][df.Exited == 1]
```

```
churn_no=df['Tenure'][df.Exited == 0]
```

```
import matplotlib.pyplot as plt
```

```
plt.xlabel("Tenure")
```

```
plt.ylabel("No of Customers")
```

```
plt.title("churn prediction ")
```

```
plt.hist([churn_yes,churn_no],color=["green","red"],label=["churn=yes","churn=no"])
```

```
plt.legend()
```

```
plt.show()
```

```
/usr/local/lib/python3.7/dist-packages/numpy/core/fromnumeric.py:3208: VisibleDeprecationWarning:
Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays
with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object'
when creating the ndarray.
```

```
return asarray(a).size
```

```
/usr/local/lib/python3.7/dist-packages/matplotlib/cbook/__init__.py:1376: VisibleDeprecationWarning:
Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays
with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object'
when creating the ndarray.
```

```
X = np.atleast_1d(X.T if isinstance(X, np.ndarray) else np.asarray(X))
```

```
sns.pairplot(data=df,hue='Exited')
```

```
<seaborn.axisgrid.PairGrid at 0x7fd191908a10>
```

```
df.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
```

```
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
```

```
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
```

```
      dtype='object')
```

```
feature=df[['CreditScore','Geography',
```

```
          'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
```

```
          'IsActiveMember', 'EstimatedSalary']]
```

```
feature
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard		
		IsActiveMember		EstimatedSalary						
0	619	France	Female	42	2	0.00	1	1	1	101348.88

1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
3	699	France	Female	39	1	0.00	2	0	0	93826.63
5	645	Spain	Male	44	8	113755.78	2	1	0	149756.71
8	501	France	Male	44	4	142051.07	2	0	1	74940.50
...
9990	714	Germany		Male	33	3	35016.60	1	1	0
	53667.08									
9995	771	France	Male	39	5	0.00	2	1	0	96270.64
9996	516	France	Male	35	10	57369.61	1	1	1	101699.77
9997	709	France	Female	36	7	0.00	1	0	1	42085.58
9998	772	Germany		Male	42	3	75075.31	2	1	0
	92888.52									

4993 rows × 10 columns

label=df['Exited']

label

0 1

1 0

3 0

5 1

8 0

..

9990 0

9995 0

9996 0

9997 1

9998 1

Name: Exited, Length: 4993, dtype: int64

```
from sklearn.compose import ColumnTransformer
```

```
from sklearn.preprocessing import OneHotEncoder
```

```
ct = ColumnTransformer([("oh",OneHotEncoder(),[1,2])],remainder="passthrough")
```

```
feature_onehot= ct.fit_transform(feature)
```

```
feature_onehot
```

```
array([[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
        1.0000000e+00, 1.0134888e+05],
       [0.0000000e+00, 0.0000000e+00, 1.0000000e+00, ..., 0.0000000e+00,
        1.0000000e+00, 1.1254258e+05],
       [1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
        0.0000000e+00, 9.3826630e+04],
       ...,
       [1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
        1.0000000e+00, 1.0169977e+05],
       [1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
        1.0000000e+00, 4.2085580e+04],
       [0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
        0.0000000e+00, 9.288520e+04]])
```

```
len(feature_onehot)
```

4993

```
feature_onehot[0]
```

```
array([1.0000000e+00, 0.0000000e+00, 0.0000000e+00, 1.0000000e+00,
```

```

0.0000000e+00, 6.1900000e+02, 4.2000000e+01, 2.0000000e+00,
0.0000000e+00, 1.0000000e+00, 1.0000000e+00, 1.0000000e+00,
1.0134888e+05])

df["Geography"].unique()

array(['France', 'Spain', 'Germany'], dtype=object)

df["Gender"].unique()

array(['Female', 'Male'], dtype=object)

from sklearn.model_selection import train_test_split

trainX,testX,trainY,testY = train_test_split(feature_onehot,label,test_size=0.2,random_state=0)

trainX

array([[0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
0.0000000e+00, 8.4300400e+04],
[0.0000000e+00, 0.0000000e+00, 1.0000000e+00, ..., 1.0000000e+00,
1.0000000e+00, 1.4203307e+05],
[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
1.0000000e+00, 1.6737626e+05],
...,
[0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
0.0000000e+00, 3.8270470e+04],
[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
0.0000000e+00, 1.1812088e+05],
[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
1.0000000e+00, 9.7755290e+04]])

testX

array([[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,

```

```
1.0000000e+00, 1.1045799e+05],  
[0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,  
0.0000000e+00, 6.3981370e+04],  
[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,  
0.0000000e+00, 1.1343608e+05],  
...,  
[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,  
0.0000000e+00, 2.6450570e+04],  
[0.0000000e+00, 0.0000000e+00, 1.0000000e+00, ..., 1.0000000e+00,  
0.0000000e+00, 5.4947510e+04],  
[0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,  
0.0000000e+00, 1.6318162e+05]])
```

trainY

3935 1

34 0

4189 0

5100 0

5918 0

..

9864 0

6541 0

3333 0

5298 0

5530 0

Name:Exited, Length: 3994, dtype: int64

```
testY
```

```
2378  0
```

```
8392  1
```

```
8410  0
```

```
4970  0
```

```
7674  0
```

```
..
```

```
7618  0
```

```
5529  0
```

```
2262  1
```

```
7122  0
```

```
7061  0
```

```
Name: Exited, Length: 999, dtype: int64
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
trainX_scale = scaler.fit_transform(trainX)
```

```
testX_scale = scaler.transform(testX)
```

```
trainX_scale
```

```
array([[ -1.0305103 ,  1.81765764, -0.58139784, ...,  0.64211021,
```

```
        -0.97918504, -0.37178651],
```

```
        [-1.0305103 , -0.55015861,  1.71999262, ...,  0.64211021,
```

```
        1.02125744,  0.97885865],
```

```
        [ 0.97039302, -0.55015861, -0.58139784, ...,  0.64211021,
```

```
        1.02125744,  1.57175791],
```

```
...,
```

```
[-1.0305103 , 1.81765764, -0.58139784, ..., 0.64211021,  
-0.97918504, -1.44864824],  
[ 0.97039302, -0.55015861, -0.58139784, ..., 0.64211021,  
-0.97918504, 0.41943738],  
[ 0.97039302, -0.55015861, -0.58139784, ..., 0.64211021,  
1.02125744, -0.05701184]])
```

testX_scale

```
array([[ 0.97039302, -0.55015861, -0.58139784, ..., 0.64211021,  
1.02125744, 0.24016548],  
[-1.0305103 , 1.81765764, -0.58139784, ..., 0.64211021,  
-0.97918504, -0.84714647],  
[ 0.97039302, -0.55015861, -0.58139784, ..., 0.64211021,  
-0.97918504, 0.30983735],  
...,  
[ 0.97039302, -0.55015861, -0.58139784, ..., 0.64211021,  
-0.97918504, -1.72517262],  
[-1.0305103 , -0.55015861, 1.71999262, ..., 0.64211021,  
-0.97918504, -1.05849196],  
[-1.0305103 , 1.81765764, -0.58139784, ..., 0.64211021,  
-0.97918504, 1.47362508]])
```

trainY

3935 1

34 0

4189 0

5100 0

5918 0

..

9864 0

6541 0

3333 0

5298 0

5530 0

Name: Exited, Length: 3994, dtype: int64

testY

2378 0

8392 1

8410 0

4970 0

7674 0

..

7618 0

5529 0

2262 1

7122 0

7061 0

Name: Exited, Length: 999, dtype: int64